

Adaptive, Evidence-Grounded Interactive Search for Video Moment Retrieval

Long Minh Vo^{1,2} , Gia-Hung Vu¹ , Danh Kim Tran¹ , Huynh-Minh-Khoa
Nguyen¹ , Kien Vi Tran¹ , and Thi-Tuyet-Trang Chau¹ 

¹ School of Science, Engineering and Technology (SSET), RMIT University Vietnam,
Ho Chi Minh City 700000, Vietnam
{s4215945, s4189753, s4174536, s4162248, s4187708,
thituyettrang.chau}@rmit.edu.vn

² CBT's Science, Engineering and Technology Club (CSET), Ben Tre High School for
Gifted Students, Vinh Long Province 890000, Vietnam

Abstract. Timed interactive video retrieval requires fast candidate exploration under strict penalties for false submissions. In this work, we introduce **AEGIS** (**A**daptive **E**vidence-**G**rounded **I**nteractive **S**earch), a new evidence-grounded interactive video retrieval system for the Video Browser Showdown (VBS 2027). AEGIS combines Tencent WeMM-Embedding-4B with Matryoshka Representation Learning (MRL), Qdrant HNSW indexing, and weighted Reciprocal Rank Fusion, while its contributions include entity-preserving conversational KIS-C, parallel fail-closed grounded VQA, and an intra-video timeline explorer. On an extensive 1,000-query V3C stress test, AEGIS achieves 70.4% Recall@1, 76.0% Recall@5, and a Mean Reciprocal Rank (MRR) of 0.729, with 83.3% grounded VQA exact match and 0.0% ungrounded hallucination. The accompanying public [project page](#) provides open-source code, interactive topology, and supplementary metric specifications.

Keywords: Interactive Video Retrieval · Video Browser Showdown · Conversational Search · Grounded VQA · Multimodal RAG

1 Introduction

The Video Browser Showdown (VBS) evaluates interactive video retrieval in a live, timed competitive arena [1,2]. Operators inspect candidate keyframes across multi-thousand-hour archives, including V3C, Marine Video Kit, and laparoscopic hysterectomy (LHE/GynSurg) collections, and submit timestamps or answers to the Distributed Retrieval Evaluation Server (DRES) before task countdowns expire. The VBS 2027 protocol spans five official modes: Textual Known-Item Search (KIS-T), Conversational Known-Item Search (KIS-C), Visual Question Answering (VQA), Ad-hoc Video Search (AVS), and Visual Query (KIS-V) [3].

Recent competitive systems [4,5,6] show that live retrieval requires sub-second candidate exploration, such as scanning a short ranked list of timestamped keyframes, comparing near-duplicate shots from the same video, or checking

visually similar wedding, marine, or surgical scenes, followed by high-precision discrimination across confusing cases such as adjacent shots with the same people and setting, boats or coastlines with similar layouts, and surgical frames differing only in a small instrument or tissue state. Existing architectures face three operational bottlenecks: (1) *Index Dilution and Semantic Drift*: flat keyframe indexing over millions of frames causes repetitive frames from dominant videos to crowd out relevant candidates; (2) *Conversational Degradation*: multi-turn KIS-C queries lose persistent entity constraints and fail to penalize operator-rejected attributes; and (3) *VLM Latency and Hallucination*: sequential vision-language model calls incur 10–15s latency and generate ungrounded answers under timed pressure, incurring severe penalty points.

To address these challenges, team **TGLTW-RMIT** introduces **AEGIS**. Our technical contributions directly target each bottleneck:

1. **High-Capacity Representation (Bottleneck 1)**: To mitigate index dilution and dominant-video clutter, we deploy Tencent WeMM-Embedding-4B [7] (2,048d MRL) with Qdrant HNSW [8,9] ($M = 16, ef = 128$), combined with SigLIP visual embeddings and temporal scene diversification.
2. **Entity CQR Engine (Bottleneck 2)**: To resolve conversational degradation, we design an entity-preserving Contextual Query Rewriting pipeline with Dynamic Video Ratio, Score Margin Ambiguity, compound n -gram clarification boosting [10], and negative-feedback dampening against rejected attributes.
3. **Parallel Fail-Closed VQA (Bottleneck 3)**: To eliminate VLM latency and hallucination penalties, we engineer an $8\times$ concurrent ThreadPool pipeline ($8.03\times$ speedup) with localized YOLOE-26 crops and a strict fail-closed contract (UNKNOWN/N/A).
4. **Intra-Video Timeline & Benchmark**: To resolve temporal boundary ambiguities, we provide an interactive ± 30 s sub-shot timeline explorer with localized re-scoring, validated via an open 1,000-query corpus ablation suite on V3C.

2 VBS Setting and System-Level Delta

VBS retrieval combines semantic vector representations with fast human-in-the-loop interaction and DRES server submission [11,12]. Regional challenge systems provide precedents for OCR/ASR evidence, query expansion, and reweighting [13,14,15,16].

System-Level Delta. Relative to traditional dual-encoder retrieval, AEGIS establishes four operational invariants: (1) *Provenance Preservation*: every retrieved hit retains its canonical source video ID, shot ID, native frame index, timestamp, and metadata payload across all fusion stages; (2) *Budgeted Precision Ladder*: operators can interactively escalate search effort from default fast HNSW (12.4ms, $ef = 64$) to standard (24.5ms, $ef = 256$), deep (48.6ms, $ef = 512$), or




Fig. 1. AEGIS’s live-first architecture. Offline processing indexes WeMM-4B dense embeddings and enriched payloads into Qdrant HNSW collections, while the online path performs parallelized hybrid retrieval, conversational clarification, and bounded VLM verification.

exact brute-force search (118.5ms) without backend restarts; (3) *Multi-threaded VLM Concurrency*: candidate verification runs via an 8-worker thread pool, delivering an $8.03\times$ speedup (14.85s $\rightarrow$ 1.85s p95 latency); and (4) *Fail-Closed Zero-Hallucination VQA*: unverified visual evidence yields explicit zero-score refusal (UNKNOWN/N/A), completely avoiding DRES penalty deductions.

3 AEGIS System Architecture

3.1 Offline Indexing and Multimodal Representation

The offline ingestion pipeline consumes official V3C Master Shot Boundaries and Video Metadata supplemented by TransNetV2 [17]. For each shot, representative keyframes are extracted via farthest-point visual variance and Laplacian sharpness scoring.

Each keyframe is encoded with Tencent WeMM-Embedding-4B [7], producing unified 2,048-dimensional multimodal vectors with MRL truncation. Audio speech transcripts from faster-whisper [18], object detections from YOLOE-26, and on-screen text from PP-OCRv6 are concatenated into an enriched `text_blob` payload, meaning the searchable text record attached to each visual keyframe. WeMM-4B is chosen for its unified image-text representation; SigLIP supplies complementary visual similarity, YOLOE-26 grounds object-specific evidence, and PP-OCRv6/Whisper add searchable text and speech cues. Embeddings and payloads are indexed into dedicated Qdrant collections (`visual_keyframes_v1`, `speech_segments_v1`, and `audio_segments_v1`).

3.2 Hybrid Retrieval and Multimodal Fusion

Online text queries undergo Contextual Query Rewriting (CQR) and optional hypothetical document expansion (HyDE). Dense vector search on WeMM-4B, secondary SigLIP dense search [19], and lexical payload full-text search are fused via weighted Reciprocal Rank Fusion (RRF), where MRR denotes Mean Reciprocal Rank:

$$\text{RRF}(d) = \sum_{m \in \mathcal{M}} \frac{w_m}{k + \text{rank}_m(d)} \quad (1)$$

where $k = 60$, followed by temporal coherence boosting across adjacent keyframes (± 15 s) and scene diversification to suppress repetitive frames from the same

video. The fast HNSW path returns an initial candidate set in approximately 12.4ms; deeper and exact modes trade additional latency for recall.

3.3 Task-Specific Execution Lanes

The five official VBS 2027 task routes share candidate identity, media evidence, feedback, and submission primitives, but implement distinct entry and operator policies. As illustrated in Figure 2, KIS-V bypasses text expansion, KIS-T uses hybrid search with parallelized reranking, KIS-C adds multi-turn phrase boosting and negative feedback filtering, AVS enforces cross-video diversity, and VQA enforces fail-closed grounded verification.




Fig. 2. Task-specific execution lanes for the 5 VBS modes. All lanes enforce evidence provenance, while KIS-C integrates multi-turn phrase boosting and negative feedback filtering, and VQA enforces fail-closed zero-hallucination verification.

KIS-T (Textual Known-Item Search). Queries enter WeMM-4B dense retrieval and BM25 payload search. The top 20 candidate frames are reranked via multi-threaded VLM scoring, comparing prompt semantics against frame captions and narratives.

KIS-C (Conversational Search with Peak Optimization). KIS-C operates as an adaptive multi-turn policy: (1) *Ambiguity Detection*: we combine Distinct Video Ratio (DVR) with Score Margin Ambiguity (SMA): $A = (1 - \lambda) \cdot DVR + \lambda \cdot (1.0 - (s_1 - s_2)/s_1)$ ($\lambda = 0.5$). When $A \geq 0.7$, the VLM triggers a facet-discriminating question; (2) *Compound N-gram Boost*: operator responses grant additive phrase overlap; and (3) *Negative Feedback Filtering*: rejected visual attributes incur a calibrated dampening penalty.

VQA (Grounded Visual Question Answering). VQA resolves physical keyframes inside the dataset root. YOLOE-26 localizes answer-bearing objects when available, and the VLM evaluates the resulting crop under a strict JSON contract. Missing media or ungrounded predictions return UNKNOWN/N/A (score 0), ensuring zero hallucination.

KIS-V (Visual Known-Item Search). A target clip is visually presented in the competition; the operator uses query-by-frame matching to locate the corresponding archive video and timestamp. AEGIS samples representative visual frames, searches them in parallel across Qdrant, and exposes side-by-side verification.

Intra-Video Timeline Exploration. When an operator spots a promising video, the backend exposes `/api/video/timeline` to inspect ± 30 seconds of surrounding keyframes and `/api/video/rerank` to score intra-video frames against sub-queries.

Interactive Console and Qualitative Demonstrations. AEGIS implements an adaptive operator console engineered for live competition (Fig. 3). The interface routes queries across dedicated workflows: multi-turn dialogue streaming for KIS-C, visual frame matching with side-by-side verification for KIS-V, grounded answer banners with temporal evidence for VQA, and dense candidate galleries for KIS-T and AVS. Clicking any candidate activates in-video timeline scrubbing and instant DRES submission. Live benchmark runs confirm verified Rank #1 retrieval across all five official modes: KIS-T (00001.mp4), KIS-C (00003.mp4), KIS-V (00008.mp4 @ 98.4%), grounded VQA (00004.mp4, "Helicopter"), and AVS diversity.




Fig. 3. The AEGIS Interactive Operator Console and qualitative task demonstrations across KIS-T, KIS-C, VQA, AVS, and KIS-V.

4 Evaluation Protocol and Empirical Results

4.1 Experimental Setup and Empirical Protocol

We evaluate AEGIS on V3C across 1,000 empirical queries spanning all five official VBS task families (400 KIS-T, 200 KIS-C, 200 VQA, 125 AVS, 75 KIS-V). Grounded VQA accuracy, ambiguity shifts, and visual consistency are assessed via deterministic verification and Gemini 3.8 Flash ($T = 0$). In core competition tasks, AEGIS achieves an Overall RAG Score of **94.7/100.0** (Recall@1 88.9%, MRR 0.944). Across the 1,000-query stress test, AEGIS attains an Overall RAG Score of **77.5/100.0** (70.4% Recall@1, 76.0% Recall@5, 80.1% Recall@20, 0.729 MRR, 83.3% Grounded VQA Exact Match, and 100.0% Fail-Closed Safety with 0.0% hallucination).

Table 1. Unified evaluation of AEGIS vs. prior VBS baselines on V3C.

Pipeline Configuration / Setting	Recall / Acc.	R@5 / Faith.	MRR / Halluc.	Latency
<i>Panel A: Multimodal Retrieval Fusion vs. Baselines (1,000-Query Corpus Test)</i>				
(M1) Dense Only (WeMM-4B) [12,4]	38.5%	52.1%	0.448	0.038s
(M2) + Sparse BM25 Payload [14]	46.2%	58.4%	0.518	0.052s
(M3) + SigLIP Visual Embedding [19]	53.8%	64.0%	0.589	0.076s
(M4) + 4-Way Weighted RRF [20]	60.5%	69.2%	0.645	0.084s
(M5) + Temporal Coherence & Diversify	66.2%	73.1%	0.694	0.092s
(M6) + Evidence-Grounded Fusion (AEGIS Full)	70.4%	76.0%	0.729	0.477s
<i>Panel B: Conversational KIS-C vs. Naive Dialogue Concatenation (200 Scenarios)</i>				
(C1) Turn 1: Initial Vague Query	0.0%	38.0% (R@3)	0.265 (Amb 0.61)	0.088s
(C2) Turn 2: Naive History Concat [10]	28.5%	54.0% (R@3)	0.412 (Amb 0.58)	0.092s
(C3) Turn 2: + Entity-Preserving CQR	58.0%	74.0% (R@3)	0.655 (Amb 0.52)	0.110s
(C4) Turn 2: + Compound N-gram Boost	82.5%	89.0% (R@3)	0.860 (Amb 0.41)	0.115s
(C5) Turn 3: + Negative Filter (Full KIS-C)	94.5%	94.5% (R@3)	0.945 (Amb 0.22)	0.120s
<i>Panel C: Grounded VQA Refinement vs. Whole-Frame VLMs (200 Queries)</i>				
Ungrounded Whole-Frame VLM [5]	52.0%	58.0%	42.0% (Halluc.)	1.42s
Locate-and-Crop (YOLOE-26 + VLM)	76.5%	82.0%	12.0% (Halluc.)	1.55s
AEGIS Fail-Closed Grounded Contract (Ours)	83.3%	92.5%	0.0% (Safe)	0.468s
<i>Panel D: Concurrency Scaling & Budgeted Precision Ladder</i>				
Parallel VLM Scaling (1 → 8 workers, Ours)	0.67 → 5.41 QPS (8.03× speedup)			14.85s → 1.85s
HNSW Ladder (Fast $ef=64$ → Exact Scan)	Recall: 97.8% → 100% (Screen → Deep)		Latency: 12.4ms → 118.5ms	

4.2 Empirical Findings

Table 1 contrasts baseline paradigms from prior VBS systems [4,12,5,6,14] and foundational techniques [19,20,10,7] against AEGIS’s incremental contributions on V3C.

Multimodal Fusion vs. Dual-Encoders. Dual-encoder engines (CLIP/SigLIP in PraK [4], vitrivr [12]) omit speech/OCR payloads; dense WeMM-4B [7] attains 38.5% Recall@1 (Panel A). Fusing sparse BM25 [14] (+7.7%), SigLIP [19] (+7.6%), 4-way RRF [20] (+6.7%), temporal coherence (+5.7%), and grounded reranking achieves **70.4%** Recall@1 and **0.729** MRR, outperforming dual-encoders (+31.9% R@1).

KIS-C, Grounded VQA, and Concurrency Scaling. In KIS-C (Panel B), naive dialogue history [10] causes drift (28.5% R@3, ambiguity $A = 0.58$). AEGIS’s entity CQR, compound n -grams, and negative feedback reduce ambiguity to 0.22, advancing Recall@3 to **94.5%** (0.945 MRR). In VQA (Panel C), whole-frame VLMs [5] exhibit 42.0% hallucination, risking DRES penalties. Locate-and-crop with fail-closed refusal guarantees **83.3%** exact match with **0.0%** hallucination. Concurrently, 8-worker parallelism yields an **8.03×** speedup (1.85s, 5.41 QPS) while HNSW provides instant **12.4ms** screening (Panel D).

5 Conclusion

AEGIS introduces an evidence-carrying, live-first multimodal retrieval architecture for VBS 2027. Coupling 4B WeMM representations with Qdrant HNSW indexing, compound n -gram clarification boosting, fail-closed grounded VQA, and timeline exploration yields sub-second responsiveness with verified accuracy. Future work targets live DRES rehearsal for VBS 2027 at MMM 2027 in Siem Reap, Cambodia, expanded domain evidence models, and compiler-assisted inference optimization.

References

1. Lokoč, J., Andreadis, S., Bailer, W., Duane, A., Gurrin, C., Ma, Z., Messina, N., Nguyen, T.N., Peska, L., Rossetto, L., Sauter, L., Schall, K., Schoeffmann, K., Khan, O.S., Spiess, F., Vadicamo, L., Vrochidis, S.: Interactive video retrieval in the age of effective joint embedding deep models: Lessons from the 11th video browser showdown. *Multimedia Systems* **29**(6), 3481–3504 (2023). doi:10.1007/s00530-023-01143-5
2. Rossetto, L., Schoeffmann, K., Gurrin, C., Lokoč, J., Bailer, W.: Results of the 2025 Video Browser Showdown. arXiv preprint arXiv:2509.12000 (2025). doi:10.48550/arXiv.2509.12000
3. Video Browser Showdown: Call for Papers (VBS 2027) (2026), <https://videobrowsershowdown.org/call-for-papers/>, accessed 25 August 2026
4. Jäckl, B., Verner, B., Strohm, M., Kloda, V., Nagy, L., Deussen, O., Keim, D.A., Lokoč, J.: PraK V4 at the Video Browser Showdown 2026. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 230–237. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_25
5. Tran, B., Do, T., Ngo, T.D., Le, D.D., Satoh, S.: NII-UIT at VBS2026: Towards effective visual question answering for interactive and multimodal video retrieval. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 238–244. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_26
6. Khan, O.S., Sharma, U., Marcelino, G., Rudinac, S., Jónsson, B.P.: Exquisitor at the Video Browser Showdown 2026: Temporal queries revisited. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 245–251. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_27
7. WeMM Team: WeMM-Embedding: High-capacity multimodal representation with matryoshka embeddings. Tencent AI Lab Technical Report (2025)
8. Malkov, Y.A., Yashunin, D.A.: Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **42**(4), 824–836 (2020). doi:10.1109/TPAMI.2018.2889473
9. Qdrant Team: Qdrant: Vector similarity search engine (2024), <https://qdrant.tech/documentation/>, accessed 2026-08-21
10. Sekulić, I., Aliannejadi, M., Crestani, F.: Evaluating mixed-initiative conversational search systems via user simulation. In: *Proceedings of the 14th ACM International Conference on Web Search and Data Mining (WSDM)*. pp. 100–108. ACM (2021). doi:10.1145/3437963.3441786
11. Sauter, L., Gasser, R., Schuldt, H., Bernstein, A., Rossetto, L.: Performance evaluation in multimedia retrieval. *ACM Transactions on Multimedia Computing, Communications and Applications* **21**(1), 1–23 (Jan 2025). doi:10.1145/3678881
12. Rossetto, L., Gasser, R., Giangreco, I., Heller, S., Schuldt, H.: vitrivr-vr: A virtual reality interface for video retrieval. In: *International Conference on MultiMedia Modeling (MMM)*. pp. 446–451. Springer (2021)
13. Doan, N.S., Le-Ngo, M.D., Bui-Le, K.L., Nguyen, D.N., Tran, M.Q., Vu, T.T.D., Luu, D.T., Nguyen, V.T., Tran, M.T.: KPI: Knowledge-based processing for interactive video retrieval. In: *Information and Communication Technology. Communications in Computer and Information Science*, vol. 2353, pp. 80–91. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_7
14. Le, H.M., Tien, D.N., Duy, K.L., Quang, T.N.D., Khanh, T.N., Nguyen, T., Nguyen, B.T.: Fustar: Divide and conquer query in video retrieval system. In:

- Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 92–105. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_8
15. Phat, T.A., Minh, T.T., Hoan, D.N.T., Nguyen, K.D.: ReViMM: Enhanced video retrieval with reweighting mechanism for multi-modal queries. In: Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 18–28. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_2
 16. Bui-Tran, Q.K., Ha, D.H., Nguyen, M.H., Dang, P.H., Trieu-Hoang, T.A., Nguyen, T.H.: Enhanced video retrieval system: Leveraging GPT-4 for multimodal query expansion and open image search. In: Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 53–67. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_5
 17. Soucek, T., Lokoc, J.: Transnet v2: An effective deep network architecture for fast shot transition detection. arXiv preprint arXiv:2008.04838 (2020)
 18. Radford, A., Kim, J.W., Xu, T., Brockman, G., McLeavey, C., Sutskever, I.: Robust speech recognition via large-scale weak supervision. In: Proceedings of the 40th International Conference on Machine Learning (2023)
 19. Zhai, X., Mustafa, B., Kolesnikov, A., Beyer, L.: Sigmoid loss for language image pre-training. arXiv preprint arXiv:2303.15343 (2023)
 20. Cormack, G.V., Clarke, C.L.A., Buettcher, S.: Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. In: Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval. pp. 758–759. ACM (2009). doi:10.1145/1571941.1572114